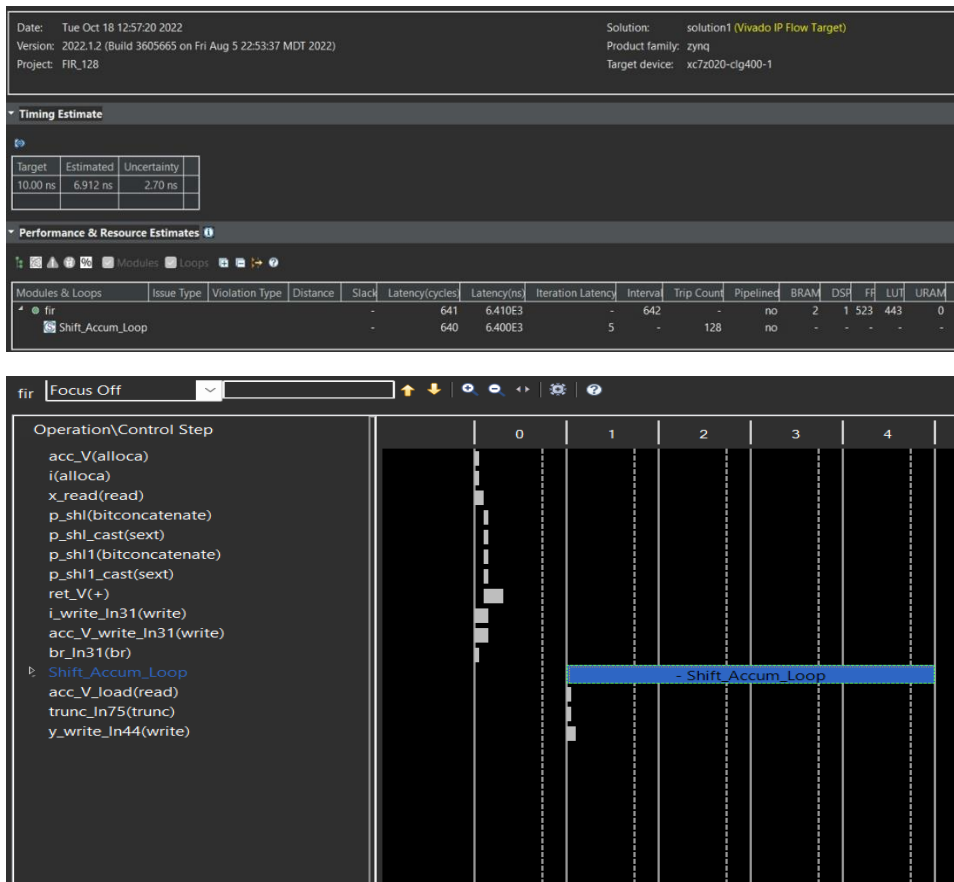


WES 237C Project 1 Report

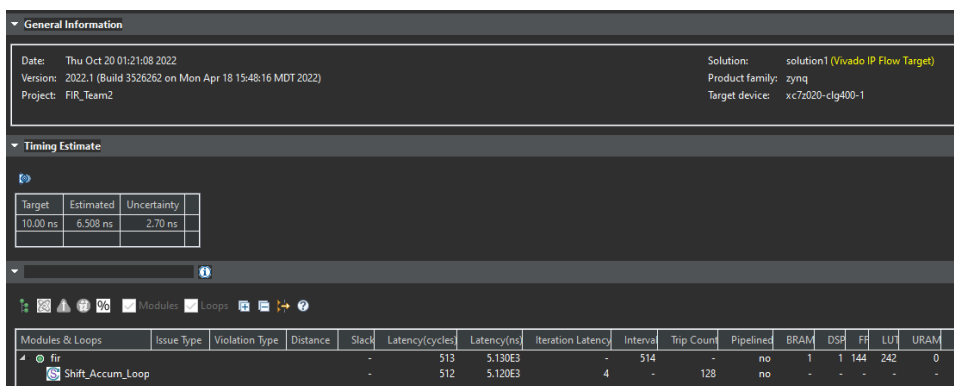
Hassan Ahmad: A59014371

Sang Ryul Pae (Eric): A59014357

* *Baseline: 6.41us*



1. Variable Bit Width: 5.12us



Using variable bit width did not show much improvement in performance. We set these following arbitrary precision data types initially:

```
#include "ap_int.h"

typedef ap_int<5>      coef_t;

typedef ap_int<17>     data_t;

typedef ap_int<22>     acc_t;
```

When we decreased the size of the largest data type, acc_t, from a bit width of 22 down to 20, we did see a slight improvement in the total number of cycles compared to the baseline performance.

2. Pipelining: 1.35us

Date: Mon Oct 17 23:19:52 2022

Version: 2022.1.2 (Build 3605665 on Fri Aug 5 22:53:37 MDT 2022)

Project: FIR_128

Solution: solution1 (Vivado IP Flow Target)

Product family: zynq

Target device: xc7z020-cig400-1

Timing Estimate

Target	Estimated	Uncertainty
10.00 ns	6.912 ns	2.70 ns

Performance & Resource Estimates

Modules Loops

Modules & Loops	Issue Type	Violation Type	Distance	Slack	Latency(cycles)	Latency(ns)	Iteration Latency	Interval	Trip Count	Pipelined	BRAM	DSH	FF	LU	URAM
• fir	-	-	-	-	135	1.350E3	-	136	-	no	3	1	755	503	0
• fir_Pipeline_Shift_Accum_Loop	-	-	-	-	133	1.330E3	-	133	-	no	3	1	576	276	0

With no II set, we got the above latency of 135 cycles. When we set II to 1, we got a low latency of 116 cycles, but a slightly longer estimated clock cycle at 7.186 ns. So, II doesn't help performance in this example as II doesn't give any benefit to other tasks. With II set to 2, we got the same performance as when II = 1, showing that adding excessive initiation interval did not improve the execution of the loop iterations.

===== pipeline/unroll/partitioning in **fir.cpp** =====

(<-- partitioning + unrolling gave big improvements then pipeline didn't give any further improvement whth them.)

```
void fir ( data_t *y, data_t x )
```

```
{
```

```
#pragma HLS PIPELINE
```

```
#pragma hls interface s_axilite port=return
```

```
coef_t c[N] = {...};
```

```
static data_t shift_reg[N];
```

```
acc_t acc;
```

```
int i;
```

```
acc = 0;
```

```
#pragma HLS ARRAY_PARTITION variable=c type=complete dim=1
```

```
Shift_Accum_Loop:
```

```
for(i=N-1;i>=0;i--){
```

```
#pragma HLS unroll
```

```
if(i == 0){
```

```
acc += x * c[0];
```

```
shift_reg[0] = x;
```

```
}
```

```
else{
```

```
shift_reg[i] = shift_reg[i-1];
```

```
acc += shift_reg[i] * c[i];
```

```
#pragma HLS ARRAY_PARTITION variable= shift_reg type=complete dim=1
```

```
}
```

```
}
```

```
*y = acc;
```

```
}
```

3. Removing Conditional Statements: 6.35us

Date: Thu Oct 20 07:20:26 2022

Version: 2022.1 (Build 3526262 on Mon Apr 18 15:48:16 MDT 2022)

Project: FIR_Team2

Solution: solution1 (Vivado IP Flow Target)

Product family: zynq

Target device: xc7z020-clg400-1

▼ Timing Estimate

Target

Estimated

Uncertainty

10.00 ns

6.912 ns

2.70 ns

▼ Performance & Resource Estimates ⓘ

*It didn't work without bit width setting with AP_int<>

BRAM/URAM were not used in all the cases tested but FF count is less in block/cycle cases and as it's because it's handling in block while complete handles each memory separately.

Date: Wed Oct 26 12:18:00 2022
Version: 3.0.12 (Build 3605665 on Fri Aug 5 22:53:37 MDT 2022)
Project: FIR_128

Solution: solution1 (**Vreado IP Flow Targeter**)
Product family: zynq
Target device: xc7z020-dsg400-1

Timing Estimate

Target	Estimated	Uncertainty
10.00 ns	6.471 ns	± 2.70 ns

Performance & Resource Estimates

Modules & Loops	Issue Type	Violation Type	Distance	Stack	Latency/cycle(s)	Latency(ns)	Iteration Latency	Interval	Trip Count	Pipeline	BRAM	DSP	FU	LUT	URAM
Shift_Accum_Loop	-	-	385	345083	-	386	-	-	128	no	0	0	247	401	0

Also, together with pipelining it gives the better performance to 60ns compared to 300ns with complete partitioning. So, proper partitioning gives optimum performance. We didn't try further as it's already good enough to target.

Date: Wed Oct 26 12:39:23 2022
Solution: solution1 (Viavio IP Flow Target)
Version: 2022.1.2 (Build 3605666 on Fri Aug 5 22:53:37 MDT 2022)
Product family: zynq
Project: FIR_128
Target device: xc7z020-clg400-1

Timing Estimate

Target	Estimated	Uncertainty
10.00 ns	6.902 ns	2.70 ns

Performance & Resource Estimates

Modules & Loops	Issue Type	Violation Type	Distance	Slack	Latency/cycle	Latency/cycle	Iteration Latency	Interval	Trips Count	Pipeline	BRAM	DSP	FF	LUT	URAM
- fir	-	-	257	2.570E3	-	258	-	-	no	0	0	2373	1051	-	0
- Shift_Accum_Loop	-	-	256	2.560E3	-	2	-	128	-	-	-	-	-	-	-

: c[] / shift_reg[] partitioning(complete), for Shift_Accum_Loop unrolling, AP_int, Pipelining

Date:Wed Oct 26 13:10:00 2022

Version:2022.1.2 (Build 3605665 on Fri Aug 5 22:53:37 MDT 2022)

Project:FIR_128

Solution:solution1 (Vivado IP Flow Target)

Product family:zynq

Target device:xc7z020-clg400-1

Timing Estimate

Target	Estimated	Uncertainty
10.00 ns	6.039 ns	2.70 ns

Performance & Resource Estimates

ModulesLoops

Modules & Loops	Issue Type	Violation Type	Distance	Slack	Latency/cycles	Latency/ns	Iteration Latency	Interval	Trip Count	Pipelined	BRAM	DSF	F#	LUT	URAM	
fir			-		6	60.000		-	7	-	no	0	0	3166	2620	0

$$\text{Throughput} = 1000 / ((6.508 \text{ ns}) * (513 \text{ cycles})) = \mathbf{0.2995 \text{ MHz}}$$

2. Pipelining:

$$\text{Throughput} = 1000 / ((6.912 \text{ ns}) * (135 \text{ cycles})) = \mathbf{1.072 \text{ MHz}}$$

3. Code Hoisting (Remove Conditional Statements):

$$\text{Throughput} = 1000 / ((6.912 \text{ ns}) * (636 \text{ cycles})) = \mathbf{0.2275 \text{ MHz}}$$

4. Loop Partitioning (Unrolling):


$$\text{Throughput} = 1000 / ((6.923 \text{ ns}) * (120 \text{ cycles})) = \mathbf{1.204 \text{ MHz}}$$

5. Memory Partitioning:

$$\text{Throughput} = 1000 / ((6.912 \text{ ns}) * (513 \text{ cycles})) = \mathbf{0.2820 \text{ MHz}}$$

6. Best Design (Memory Partitioning[block:32] + loop unrolling + Pipelining + bit width):

$$\text{Throughput} = 1000 / ((6.039 \text{ ns}) * (6 \text{ cycles})) = \mathbf{27.598 \text{ MHz}}$$



FIR Area Results for Seven Architectures

